Figure 1: ASVspoof 2019 Countermeasure Decision Matrix & Protocol Architecture

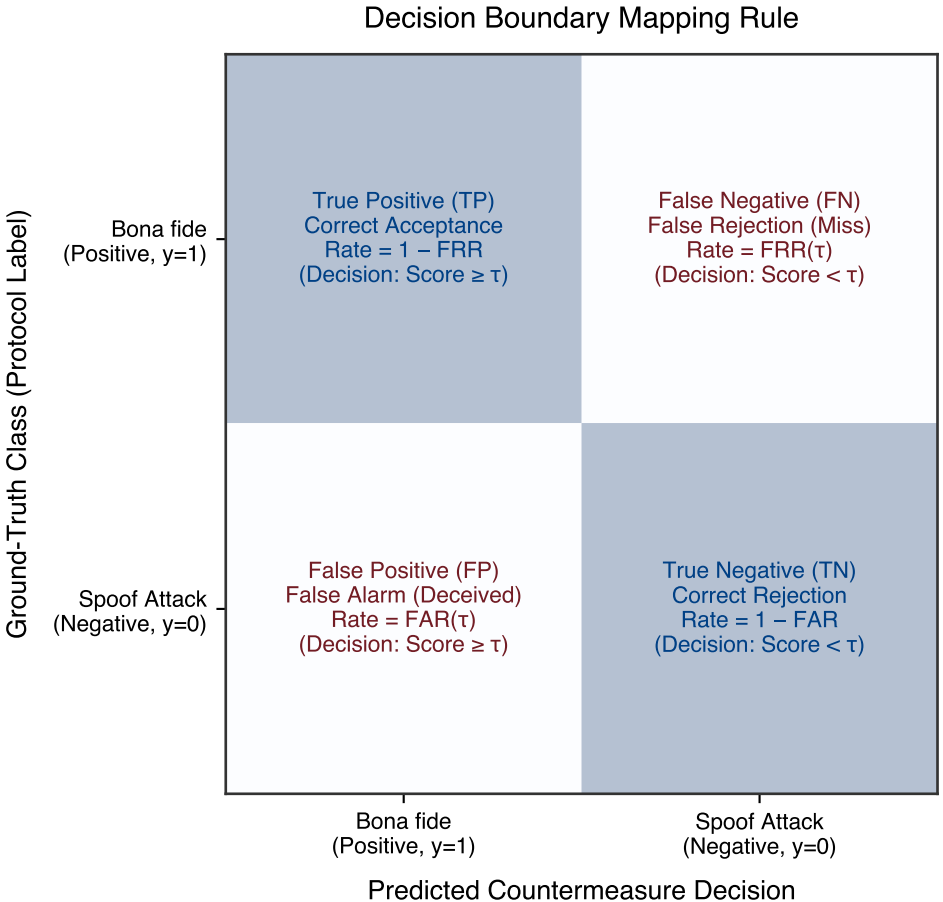


```
EVALUATION STATUS & BENCHMARK AUDIT (ADR-002)
---------------------------------------------

• Target Model Architecture: AASIST-L (85,306 Parameters)
• Scoring Convention: S_cm = Bonafide Logit (Index 1)
• Decision Rule: Predict Bonafide if S_cm ≥ τ, else Spoof
• Threshold Policy: τ is tuned on Dev EER, frozen for Eval

Official ASVspoof 2019 LA Dataset Partitions:
  1. Development Partition (Dev):
     - Total Trials: 24,844
     - Bona fide: 2,548 trials (8 speakers)
     - Spoof (Attacks A01–A06): 22,296 trials
     - Published Baseline Target: EER = 0.99%, min-tDCF = 0.0345
  2. Held-Out Evaluation Partition (Eval):
     - Total Trials: 71,237 (unseen attacks A07–A19)
     - Bona fide: 7,355 trials | Spoof: 63,882 trials

DATA STORAGE STATUS: BLOCKED_PENDING_PHYSICAL_DATASET
In accordance with ADR-002, the ~35 GiB uncompressed corpus is
held unmounted on internal SSD to prevent system volume exhaustion.
Full empirical trial confusion counts will unlock upon mounting
external media (≥ 50 GiB) and running scripts/evaluate_scores.py.
```