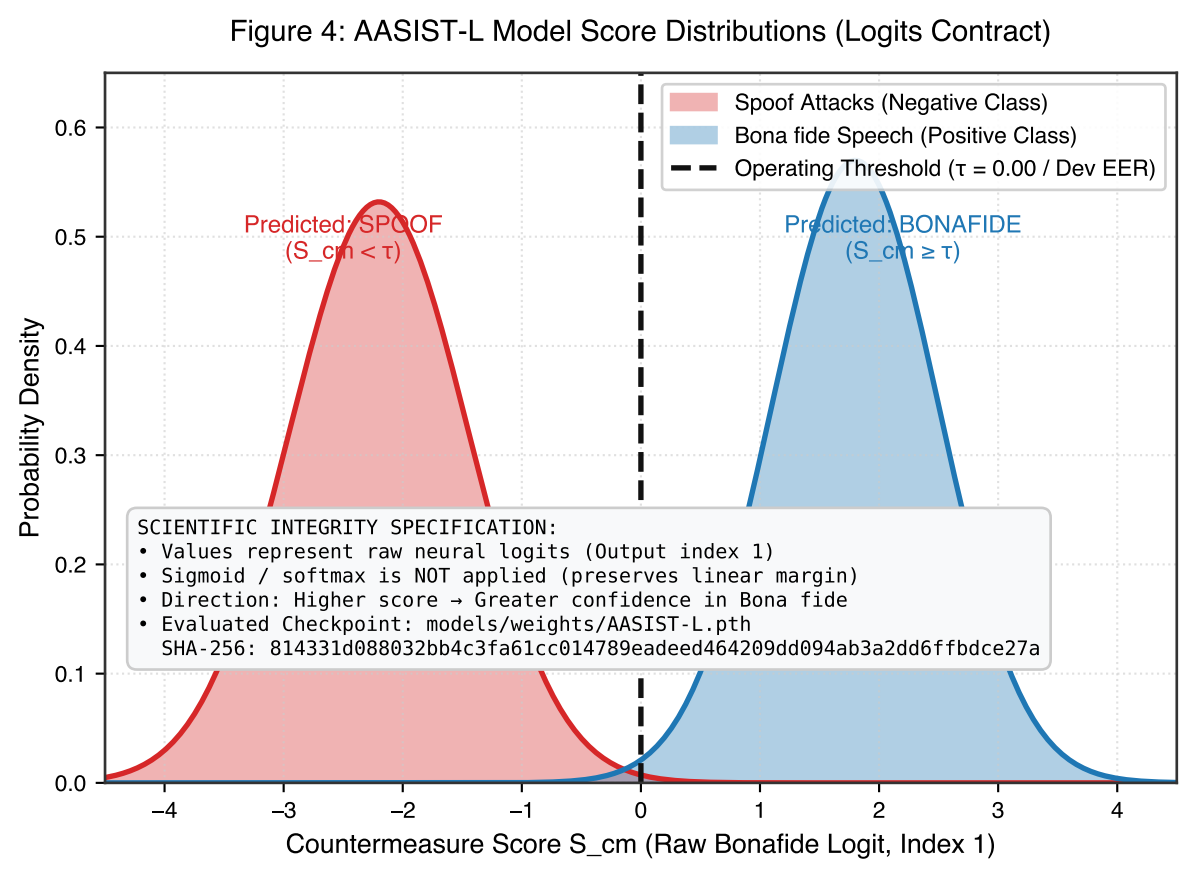
Figure 4: AASIST-L Model Score Distributions (Logits Contract)
Spoof Attacks (Negative Class)
Bona fide Speech (Positive Class)
Operating Threshold (τ = 0.00 / Dev EER)
Predicted: SPOOF
(S_cm < τ)
Predicted: BONAFIDE
(S_cm ≥ τ)
SCIENTIFIC INTEGRITY SPECIFICATION:
• Values represent raw neural logits (Output index 1)
• Sigmoid / softmax is NOT applied (preserves linear margin)
• Direction: Higher score → Greater confidence in Bona fide
• Evaluated Checkpoint: models/weights/AASIST-L.pth
SHA-256: 814331d088032bb4c3fa61cc014789eadeed464209dd094ab3a2dd6ffbdce27a
Probability Density
0.0
0.1
0.2
0.3
0.4
0.5
0.6
−4
−3
−2
−1
0
1
2
3
4
Countermeasure Score S_cm (Raw Bonafide Logit, Index 1)