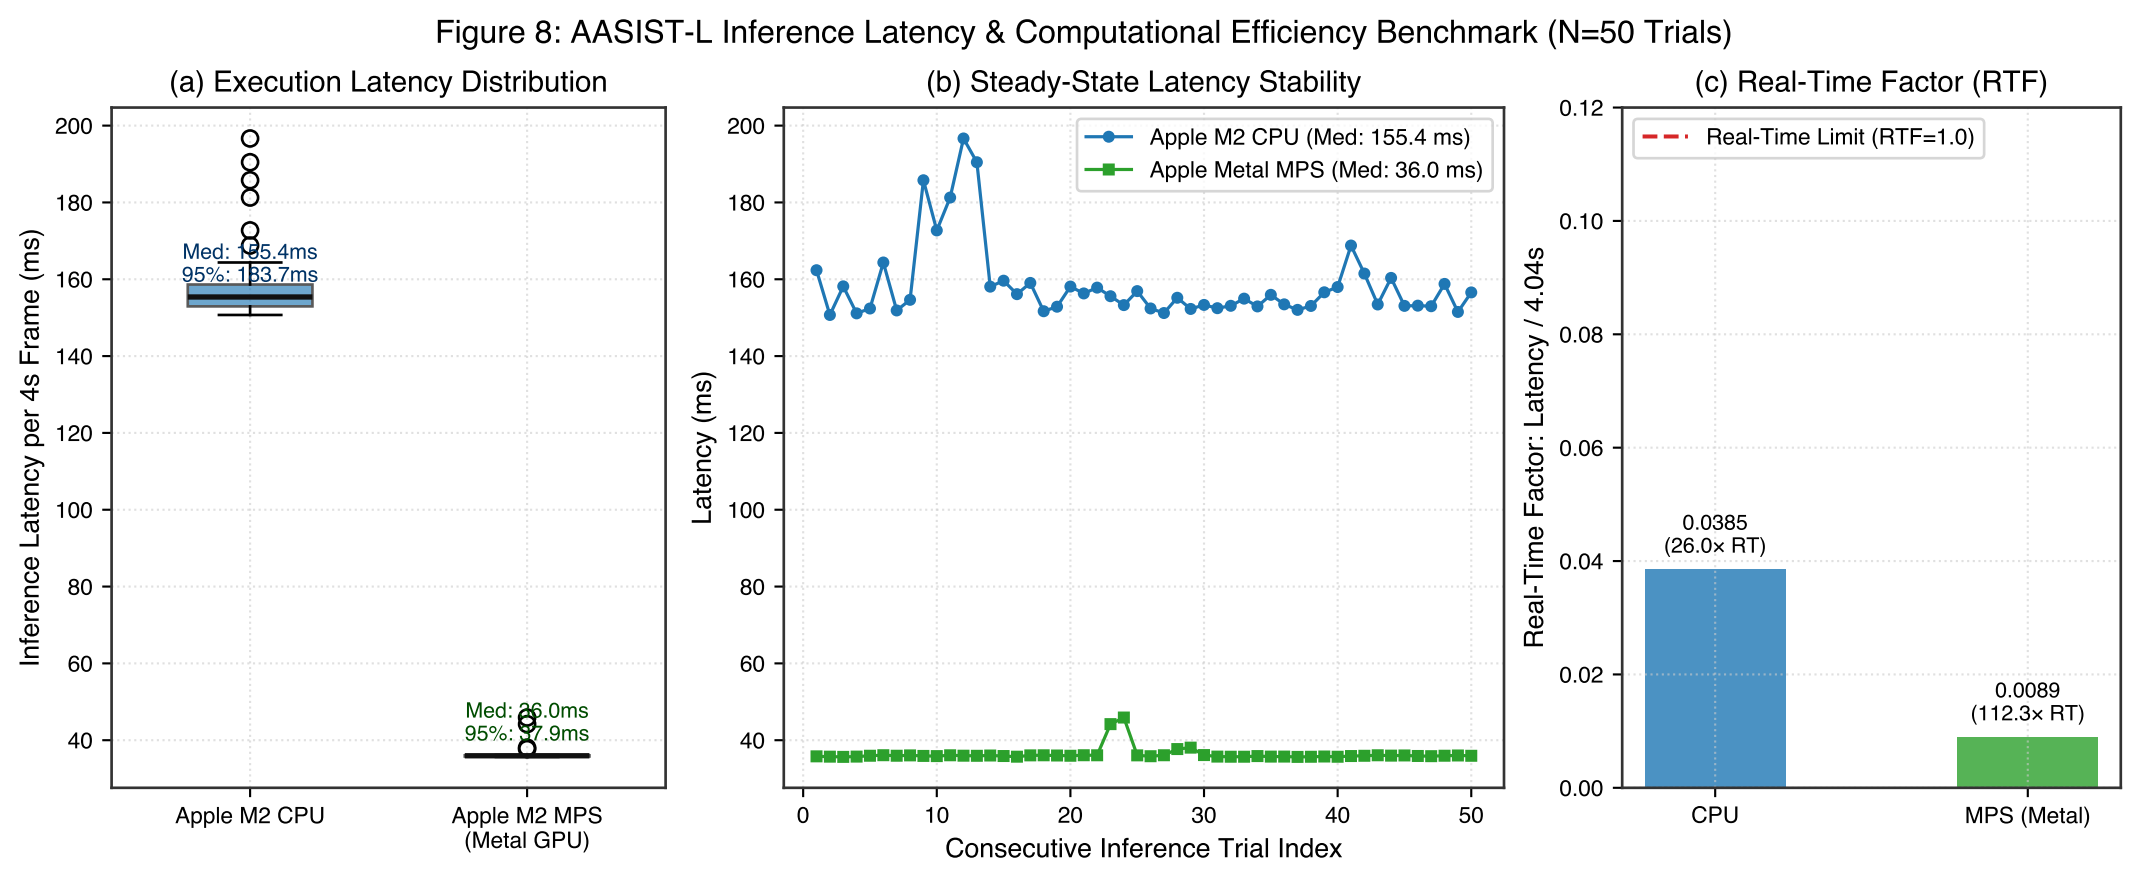
Figure 8: AASIST-L Inference Latency & Computational Efficiency Benchmark (N=50 Trials)
(a) Execution Latency Distribution
Inference Latency per 4s Frame (ms)
Med: 155.4ms
95%: 183.7ms
Med: 36.0ms
95%: 37.9ms
Apple M2 CPU
Apple M2 MPS
(Metal GPU)
(b) Steady-State Latency Stability
Latency (ms)
Apple M2 CPU (Med: 155.4 ms)
Apple Metal MPS (Med: 36.0 ms)
Consecutive Inference Trial Index
(c) Real-Time Factor (RTF)
Real-Time Factor: Latency / 4.04s
Real-Time Limit (RTF=1.0)
0.0385
(26.0× RT)
0.0089
(112.3× RT)
CPU
MPS (Metal)